

# Lab 7

Juan D Astudillo

11:59PM March 31, 2019

Generate  $\mathbb{D}$  with  $n = 100$  and  $p = 1$  where  $x$  is created from iid realizations from a standard uniform,  $y$  comes from  $f(x) = 3 - 4x$  and  $\delta$  are iid realizations from a T distribution with 10 degrees of freedom.

```
set.seed(1997)
n = 100
p = 1
X = matrix(runif(n) , ncol = 1)
f_x = 3 - 4 * X
y = f_x + rt(n, df = 10)
```

Run the linear model using `lm` and compute `b`, RMSE and  $R^2$ .

```
linear_mod = lm(y ~ X)
coef(linear_mod)
```

```
## (Intercept)          X
##      2.855769      -3.855183
```

```
summary(linear_mod)$sigma
```

```
## [1] 1.308767
```

```
summary(linear_mod)$r.squared
```

```
## [1] 0.4044066
```

Progressively add columns of  $x$  (as draws from a standard uniform), run the linear model, and show  $R^2$  goes to 1 and  $s_e$  goes to zero. Save the  $s_e$  in a vector called `in_sample_s_e`.

```
in_sample_s_e = array(NA, n - 2)
linear_mods = list()

for (j in 1 : (n - 2)){
  X = cbind(X, runif(n))
  linear_mods[[j]] = lm(y ~ ., data.frame(X))
  in_sample_s_e[j] = sd(linear_mods[[j]]$residuals)
}
dim(X)
```

```
## [1] 100  99
```

```
summary(linear_mods[[j]])$r.squared
```

```
## [1] 1
```

```
in_sample_s_e
```

```
## [1] 1.28649282 1.28232218 1.27201474 1.25707441 1.23002701 1.22829460
## [7] 1.22670007 1.22643605 1.21086228 1.16906116 1.16899113 1.14450895
## [13] 1.14430182 1.14429580 1.12528678 1.12513648 1.11418967 1.11103788
## [19] 1.10707910 1.10617314 1.10593501 1.10383989 1.09059722 1.05983015
## [25] 1.05668551 1.04443529 1.03545946 1.03542747 1.01766877 1.00972008
## [31] 0.94628966 0.94173458 0.93946800 0.92946819 0.91433407 0.91400464
```

```
## [37] 0.91178872 0.90295678 0.90275680 0.88591059 0.87734241 0.87731196
## [43] 0.87038874 0.87012819 0.86867611 0.85599859 0.84341701 0.84331509
## [49] 0.84231901 0.83898042 0.83429865 0.83191842 0.79093216 0.78712763
## [55] 0.76614183 0.74416204 0.74347964 0.74237630 0.74050592 0.72910998
## [61] 0.72569710 0.71281965 0.71163428 0.70538923 0.70016453 0.58618111
## [67] 0.58106388 0.56639234 0.55400937 0.55397187 0.54946747 0.54303693
## [73] 0.54086858 0.53230397 0.52439764 0.52245609 0.52184794 0.51818504
## [79] 0.51403634 0.51239241 0.48767310 0.46032595 0.44978526 0.32203624
## [85] 0.31974154 0.31231435 0.31111980 0.30818469 0.30506651 0.27027212
## [91] 0.22558040 0.21105584 0.19626704 0.18477877 0.17373682 0.14093467
## [97] 0.03613263 0.00000000
```

```
d = diff(in_sample_s_e)
all(d < 0)
```

```
## [1] TRUE
```

Compute a corresponding vector `oos_s_e` and show that it is increasing (for the most part) in degrees of freedom.

```
n_star = 1e5
p = 1
X_star = matrix(runif(n_star) , ncol = 1)
f_x_star = 3 - 4 * X_star
y_star = f_x_star + rt(n_star, df = 10)

oos_s_e = array(NA, n - 2)

for (j in 1 : (n - 2)){
  X_star = cbind(X_star, runif(n_star))
  y_hat_star = predict(linear_mods[[j]], data.frame(X_star))
  oos_s_e[j] = sd(y_star - y_hat_star)
}
```

```
oos_s_e
```

```
## [1] 1.138872 1.145037 1.158311 1.174007 1.203102 1.205684 1.209192
## [8] 1.211103 1.245249 1.301276 1.300797 1.306607 1.306549 1.306628
## [15] 1.333529 1.332744 1.363185 1.365351 1.371233 1.370490 1.372865
## [22] 1.375974 1.379607 1.408161 1.394222 1.426631 1.402238 1.399934
## [29] 1.438706 1.435763 1.546897 1.543536 1.543032 1.508717 1.537036
## [36] 1.538435 1.543836 1.536747 1.536909 1.576790 1.566087 1.565106
## [43] 1.593066 1.593993 1.599078 1.626604 1.616842 1.617316 1.625600
## [50] 1.618205 1.614296 1.593533 1.691716 1.717884 1.821657 1.796961
## [57] 1.787966 1.817988 1.836788 1.852059 1.863737 1.838741 1.846605
## [64] 1.912911 1.931040 2.144888 2.156903 2.236306 2.314845 2.314554
## [71] 2.324415 2.364255 2.376579 2.488367 2.417481 2.491206 2.500137
## [78] 2.420275 2.381528 2.395049 2.737752 2.694287 2.893602 3.232287
## [85] 3.276438 3.303585 3.417506 3.397106 3.440856 3.663598 3.941270
## [92] 4.244160 4.170434 4.293145 4.653674 5.214444 10.300784 14.797485
```

```
d = diff(oos_s_e)
all(d > 0)
```

```
## [1] FALSE
```

Validate the linear model for the Boston housing data.

```

Xy = MASS::Boston
K = 10
test_indices = sample(1 : nrow(Xy), 1 / K * nrow(Xy))
train_indices = setdiff(1 : nrow(Xy), test_indices)

Xy_train = Xy[train_indices, ]
Xy_test = Xy[test_indices, ]

lin_mod = lm(medv ~ ., Xy_train)
lin_mod

```

```

##
## Call:
## lm(formula = medv ~ ., data = Xy_train)
##
## Coefficients:
## (Intercept)      crim          zn          indus          chas
##  35.679799   -0.105926   0.044428   0.036199   1.785549
##          nox          rm          age          dis          rad
## -18.514293   4.075433  -0.000166  -1.454425   0.310478
##          tax      ptratio          black          lstat
##  -0.012749  -0.989053   0.009001  -0.492947

```

```
sd(lin_mod$residuals)
```

```
## [1] 4.725718
```

```

y_hat_test = predict(lin_mod, Xy_test)
sd(Xy_test$medv - y_hat_test)

```

```
## [1] 4.344533
```

```
dim(Xy)
```

```
## [1] 506 14
```

Let  $x$  be iid realizations from a  $U(0, 5)$ ,  $y$  comes from  $f(x) = 3 - 4x + 2x^2$  and  $\epsilon$  are iid realizations from a standard normal distribution. With no limit on the number of samples you can take, use regular OLS *without a quadratic term*, find the true  $h^*(x)$  (there will be no sampling variability at  $n \rightarrow \infty$  and find the oos variance of the residuals.

```

set.seed(89)
beta0 = 1
beta1 = 1
x = as.matrix(cbind(runif(n, 0, 5)))
f_x = 3 - 4*x + 2 * x^2
y = f_x + runif(n)
h_star = beta0 + beta1 * x
y_2 = 3 - 4*x + runif(n)
ols = lm(y_2 ~ x)
K = 7
test_indices = sample(1 : n, 1 / K * n)
train_indices = setdiff(1 : n, test_indices)
X_train = x[train_indices, ]
y_train = y[train_indices]
X_test = x[test_indices, ]
y_test = y[test_indices]

```

```
ols2 = lm(y_train ~ ., data.frame(X_train))
summary(ols2)$r.squared
```

```
## [1] 0.6106033
```

```
sd(ols2$residuals)
```

```
## [1] 154.145
```

```
y_hat_oos = predict(ols2, data.frame(X_test))
```

```
## Warning: 'newdata' had 14 rows but variables found have 86 rows
```

```
oos_residuals = y_test - y_hat_oos
```

```
## Warning in y_test - y_hat_oos: longer object length is not a multiple of
## shorter object length
```

```
sd(oos_residuals)
```

```
## [1] 231.9101
```

Was there any overfitting in the previous exercise? in the exercise where we used Boston data there is overfitting since the RMSE approaches 0, which is “no error”. With the oos RMSE, happens the opposite instead of going to 0, it gets larger and larger.

Find the error due to misspecification and due to ignorance expressed as variance of components of the residuals.

*#TO-DO*

At  $n = 100$ , find the error due to estimation, due to misspecification and due to ignorance expressed as variance of components of the residuals.

*#TO-DO*

Do the variances add up to the total variance of the residual?

*#TO-DO*

Validate the linear model for the Boston housing data where each feature is also modeled with a squared feature.

```
X = MASS::Boston
y = X$medv
X$medv = NULL
X = cbind(X, X^2)
colnames(X)[14 : 26] = paste(colnames(X)[1 : 13], "_sq", sep = "")
X$chas_sq = NULL
```

```
K = 10
test_indices = sample(1 : nrow(Xy), 1 / K * nrow(Xy))
train_indices = setdiff(1 : nrow(Xy), test_indices)
```

```
X_train = X[train_indices, ]
y_train = y[train_indices]
X_test = X[test_indices, ]
y_test = y[test_indices]
```

```
lin_mod = lm(y_train ~ ., X_train)
```

```
#lin_mod
sd(lin_mod$residuals)
```

```
## [1] 3.832246
```

```
y_hat_test = predict(lin_mod, X_test)
sd(y_test - y_hat_test)
```

```
## [1] 3.432411
```

Validate the linear model for the Boston housing data where each feature is also modeled with a squared feature and a cubed feature.

```
X = MASS::Boston
```

```
y = X$medv
X$medv = NULL
X = cbind(X, X^2, X^3)
```

```
colnames(X)[14 : 26] = paste(colnames(X)[1 : 13], "_sq", sep = "")
colnames(X)[27 : 39] = paste(colnames(X)[1 : 13], "_cb", sep = "")
X$chas_sq = NULL
X$chas_cb = NULL
```

```
K = 10
test_indices = sample(1 : nrow(Xy), 1 / K * nrow(Xy))
train_indices = setdiff(1 : nrow(Xy), test_indices)
X_train = X[train_indices, ]
y_train = y[train_indices]
X_test = X[test_indices, ]
y_test = y[test_indices]
```

```
lin_mod = lm(y_train ~ ., X_train)
```

```
sd(lin_mod$residuals)
```

```
## [1] 3.671813
```

```
y_hat_test = predict(lin_mod, X_test)
sd(y_test - y_hat_test)
```

```
## [1] 3.303085
```

Validate the linear model for the Boston housing data where each feature is also modeled with a squared feature and a cubed feature and a  $\log(x + 1)$  feature and an exponential feature.

```
X = MASS::Boston
```

```
y = X$medv
X$medv = NULL
X = cbind(X, X^2, X^3, log(X + 1))
```

```
colnames(X)[14 : 26] = paste(colnames(X)[1 : 13], "_sq", sep = "")
colnames(X)[27 : 39] = paste(colnames(X)[1 : 13], "_cb", sep = "")
colnames(X)[40 : 52] = paste(colnames(X)[1 : 13], "_log", sep = "")
X$chas_sq = NULL
X$chas_cb = NULL
X$chas_log = NULL
```

```

K = 10
test_indices = sample(1 : nrow(Xy), 1 / K * nrow(Xy))
train_indices = setdiff(1 : nrow(Xy), test_indices)
X_train = X[train_indices, ]
y_train = y[train_indices]
X_test = X[test_indices, ]
y_test = y[test_indices]

lin_mod = lm(y_train ~ ., X_train)

sd(lin_mod$residuals)

```

```
## [1] 3.383186
```

```

y_hat_test = predict(lin_mod, X_test)
sd(y_test - y_hat_test)

```

```
## [1] 3.434556
```

Why do we need to log  $x + 1$ ? Why not use  $\log(x)$ ?

Because if we take the log of 0 is undefined, in this way we make all the zeros 1,  $\log$  of 1 = 0.